

# PicoClaw Next-Gen

## (Codename: "AetherClaw")

Ultra-lean, sub-5 MB RAM, <500 ms boot, privacy-first, monetizable edge AI assistant. Runs on \$5 hardware, beats every Claw-family fork on adoption/retention metrics.

### 1. Structured Teardown of the Original Product (PicoClaw Go + picoclaw-rs Gateway)

Architecture (inferred from repo tree, docs, commits, Rust port mirror)

- **Single static Go binary** (Cobra CLI + pkg/ modular core).
- **Core loop**: Agent ReAct-style (LLM call → tool → iterate  $\leq 20$  times).
- **LLM layer**: `model_list` router (vendor/model prefix, OpenAI-compat + Anthropic/Zhipu special paths, round-robin load-balance, zero-code addition).
- **Tools layer** (pkg/tools inferred): `read/write/edit/append/list_dir + exec` (path + dangerous-pattern guards), `web_search` (Brave/Tavily/DuckDuckGo fallback), `spawn` (async sub-agent via message bus), `message, cron`. Sandbox flag `restrict_to_workspace` applied uniformly to main/sub/heartbeat.
- **Channels/Gateway**: Base allowlist + per-channel (Telegram long-poll via teloxide-like, webhooks for LINE/WeCom/Discord/QQ/DingTalk). picoclaw-rs reimplements Tokio/Axum/mpsc bus + health server pattern.
- **Memory/Persistence**: File-based workspace (`~/.picoclaw/workspace/` with AGENTS.md, HEARTBEAT.md, IDENTITY.md, SOUL.md, MEMORY.md, sessions/, cron/). Heartbeat every 30 min reads MD → executes or spawns.
- **Security**: Consistent boundary (no bypass via subagents), exec blocklist.
- **CLI**: `onboard / agent / gateway / cron / status`. Docker Compose profiles.
- **Boot/Perf**: Static binary → 1 s cold start, 10–20 MB RSS (recent PR bloat noted), cross-compile (x86/ARM/RISC-V).
- **Deps** (inferred minimal): Standard library heavy, Cobra, reqwest-like HTTP, JSON/MD parsers. No heavy frameworks.



## Value Delivery

- **Functional:** Personal always-on agent on \$10 LicheeRV-Nano / old Android (Termux).
- **Emotional:** “Second life for decade-old devices”, “AI that fits in your pocket”.
- **Social:** AetherClaw network, massive community PRs (12 k stars in 1 week).

## Friction & Bloat

- Config hell (deep JSON + env overrides).
- No local/offline model (cloud-only → cost + latency + privacy).
- Pure CLI + MD files = high activation barrier for non-technical users.
- Memory creep from PRs (10→20 MB).
- No built-in analytics/monetization hooks.
- Channel setup “medium” for most except Telegram.

## Benchmarks (public + inferred)

- RAM: 10–20 MB vs OpenClaw >1 GB, NanoBot >100 MB.
- Boot: <1 s on 0.6 GHz vs 30–500 s.
- Cost-to-serve: near-zero (edge).
- Reliability: Good (sandbox), but no observability/SRE.

## 2. Lean Rebuild Blueprint (Replicate → Simplify → Optimize)

Target Stack (Rust 2026 edition – beats Go on binary size + safety)

### Language

Rust (Tokio + Axum + Serde + Reqwest + tracing). Use `musl` target → static <8 MB binary.

**Minimal MVP parity** (2–3 weeks for solo dev + community)



- `config/` → `aetherclaw.toml` (or JSON) with layered defaults + auto-migrate from PicoClaw.
- `core/agent/` → single `AgentLoop` + ReAct engine (max 15 iterations default).
- `core/tools/` → trait-based, sandbox via `cap-std` + path canonicalization (faster than string guards).
- `core/l1m/` → `ModelRouter` (same `vendor/model` syntax, built-in tiny local via `llama-cpp-rs` or `candle` quantized 1–3B models).
- `channels/` → trait + registry (Telegram first via teloxide, Axum webhooks).
- `bus/` → Tokio mpsc (100 capacity) + broadcast for sub-agents.
- `heartbeat/` + `cron/` → single Tokio task with `tokio-cron-scheduler`.
- `workspace/` → same MD files, but with optional SQLite for sessions/cron (opt-in, <1 MB).

## Remove bloat

No Cobra (clap derive), no full Docker Compose (single binary `aetherclaw up`). Drop deprecated `providers` path entirely.

## Foundation = FEW

- Fast: 300–500 ms boot (Rust cold start + lazy init).
- Efficient: Target <5 MB RSS (cap-std, no GC pauses, arena alloc where possible).
- Workhorse: `tracing` + OpenTelemetry export, health `/live /ready /metrics` (Prometheus), graceful shutdown.

## Build & Release

```
cargo build --target x86_64-unknown-linux-musl --release
# <8 MB binary, works on $5 RISC-V
```

## 3. 10× Enhancement Plan (UX/UI, Speed, Affordability, Retention Loops)



## UX/UI – Frictionless Activation

- `aetherclaw init` → interactive TUI (ratatui) wizard (pick hardware, LLM, channels).
- Optional tiny Axum + HTMX web dashboard (localhost:8080) for non-devs.
- Markdown prompts auto-validated + templates gallery.
- Voice-first: built-in Whisper-tiny (local) + TTS (piper-rs).

## Performance 10×

- Local-first: 1–3B quantized models run on-device (0.5–2 tokens/s on \$10 board).
- Hybrid routing: cheap cloud fallback only when local confidence is low.
- Sub-agents: true zero-copy via bus, max 5 concurrent.
- Edge compute: optional WASM plugins for skills (sandboxed).

## Affordability

- Core = free forever (MIT).
- Local models = zero API cost.
- Optional “AetherHub” (hosted relay) at \$3/mo for unlimited premium models + backup.

## Retention & Habit Loops

- Daily “Soul Check-in” summary pushed to all channels.
- Streaks + memory milestones (“You and AetherClaw have planned 42 tasks together”).
- Auto-spawn “Life Coach” sub-agent every Sunday.
- AetherClaw federation + public skill marketplace (share/reuse HEARTBEAT.md snippets).

## Monetization Levers

- Freemium: free local + limited cloud; Pro \$5/mo (unlimited cloud + priority support + private skills).
- Marketplace: 30 % cut on paid skills (voice packs, domain-specific agents).
- API metering: sell gateway-as-service for enterprises.
- Hardware bundle with Sipeed (affiliate).



## 4. Marketability + Viability + Feasibility Analysis

### Desirability

Extremely high. Edge AI + privacy + \$5–10 hardware wave (RISC-V explosion) + “revive old Android” narrative. Claw family already proves demand (PicoClaw 19 k stars in weeks).

### Feasibility

High. Rust ecosystem mature (llama-cpp-rs, teloxide, axum all production). Start with 1–2 full-time + 10 community maintainers (exactly PicoClaw’s call). MVP in 4–6 weeks.

### Viability

Strong.

- **Business models:** Open-core + Pro subscription + marketplace (proven: Cursor, Windsurf, local LLM tools).
- **Cost-to-serve:** Near zero (edge) → 90 %+ margin on Pro.
- **Pricing:** Beats every competitor (OpenClaw/Mac-mini class = \$599+ hardware + cloud bills).
- **Go-to-market:**
  - Day 1: “PicoClaw → AetherClaw migration script” (one-liner).
  - Viral: “Runs on your \$9 LicheeRV-Nano” videos + Termux one-liner.
  - Community: Discord + WeChat + AetherClaw native.

## 5. Defensibility & Moat Design

### Proprietary Moat

- **Ultra-sandbox + cap-std + WASM skills** = auditable, verifiable security (Rust memory safety + formal-ish bounds).
- **File-based personality system + marketplace** = network effect (users share SOUL.md snippets).
- **Hybrid local/cloud router with confidence scoring** = best-of-both (privacy + power).
- **Hardware co-design hooks** (I2C/SPI tools already in PicoClaw) → official Sipeed/MaixCAM bundles.



## Open-Core Strategy

- Core engine MIT.
- AetherHub relay + premium skills = closed-source revenue.
- API-first for enterprise integrations.

## Scaling from Day 1

- Observability: OpenTelemetry → Grafana Cloud free tier.
- Global: Multi-region edge relays (Fly.io/Cloudflare Workers).
- SRE: Auto-scaling sub-agents, circuit breakers on LLM calls.
- Data advantage: Opt-in anonymized workspace insights (never raw user data) for model fine-tuning.

## 6. Scale-Ready Roadmap (Prioritized, Executable, 90-Day MVP)

### Phase 0 – Foundation (Week 1–2)

- Fork picoclaw-rs as base.
- Port all tools + sandbox (cap-std).
- Implement `model_list` router + local candle/llama-cpp 1B fallback.
- TUI init wizard.
- Migration command from PicoClaw workspace.

### Phase 1 – Parity + 5× Perf (Week 3–5)

- Full channel parity + Axum dashboard.
- Heartbeat + cron + spawn (zero-copy).
- Metrics + tracing.
- Release v0.1 “AetherClaw Beta” – target <5 MB RSS, <500 ms boot.

### Phase 2 – 10× UX & Retention (Week 6–8)

- Voice (local Whisper + TTS).
- Skill marketplace stub (Git + AetherClaw).
- Daily summary + streak system.
- AetherHub Pro backend (Stripe).



## Phase 3 – Monetization & Scale (Week 9–12)

- Marketplace launch.
- Hardware bundles (Sipeed).
- Enterprise self-hosted + relay.
- Community maintainers program (same as PicoClaw but with revenue share).

### Success Metrics

- 10 k stars in 30 days.
- <5 MB RSS on LicheeRV-Nano.
- 30 % of users on Pro within 90 days.
- 5× lower activation time vs PicoClaw (measured by init-to-first-tool-call).

### One-Line Execution Command (Day 1 for any contributor)

```
git clone https://github.com/you/aetherclaw && cd aetherclaw &&  
./bootstrap.sh
```

This plan is directly executable today. It takes everything PicoClaw proved (extreme lightness + community love) and fixes every friction while adding retention/monetization moats that turn users into lifelong advocates and paying customers.

AetherClaw isn't just faster or cheaper — it is the final form of the Claw family: the AI that disappears into your hardware and becomes part of your life.

Ready to ship? Let's build it. 皮皮虾, 我们走 ! 🦀→🚀



# AetherClaw Project: Step-by-Step Execution Guide & Public Strategy

## AetherClaw Project: Step-by-Step Execution Guide

This guide is structured into four main phases, moving from initial foundation and parity to feature enhancement, monetization, and scale.

### Phase 0: The Foundation (Weeks 1–2)

**Objective: Establish the core Rust architecture, security sandbox, and basic LLM routing.**

Step	Action Item	Details & Goal	Source/Target Metric
0.1	<b>Fork and Stabilize Base</b>	Fork <a href="#">picoclaw-rs</a> as the new base. Ensure a clean, minimal initial build.	Fork completed. <a href="#">cargo build</a> successful.
0.2	<b>Port Core Tools &amp; Sandbox</b>	Port all original <a href="#">pkg/tools</a> (read/write/exec, etc.). Implement the new security model using <b>cap-std</b> and path canonicalization for faster, safer	Tools ported. Sandbox confirmed functional.



		sandboxing.	
0.3	<b>LLM Router Implementation</b>	Implement the <b>ModelRouter</b> with the same vendor/model syntax. Add support for the tiny local fallback model via <code>candle</code> or <code>llama-cpp-rs</code> (1B quantized models).	Hybrid routing operational. Local model loads.
0.4	<b>Frictionless Initialization</b>	Build the TUI ( <code>ratatui</code> ) init wizard ( <code>aetherclaw init</code> ) to guide users through hardware, LLM, and channel setup.	TUI wizard completed.
0.5	<b>PicoClaw Migration</b>	Implement the one-line migration command to auto-migrate workspaces from PicoClaw (JSON config and MD files) to <code>aetherclaw.toml</code> .	Migration script validated.



## Phase 1: Parity & Performance (Weeks 3–5)

**Objective: Achieve functional parity with the original PicoClaw while hitting key performance targets.**

Step	Action Item	Details & Goal	Source/Target Metric
1.1	<b>Channel Parity &amp; Dashboard</b>	Implement full channel parity (Telegram first). Launch the minimal <b>Axum + HTMX web dashboard</b> for non-dev users (localhost:8080).	Full channel support. Web dashboard functional.
1.2	<b>Heartbeat &amp; Cron Tasks</b>	Implement the <b>Heartbeat</b> and <b>cron</b> functionality using a single Tokio task with <b>tokio-cron-scheduler</b> . Ensure the <b>spawn</b> sub-agent process uses a <b>zero-copy</b> message bus.	Heartbeat/Cron active. Zero-copy sub-agents confirmed.
1.3	<b>Observability</b>	Integrate <b>tracing</b> and <b>OpenTelemetry</b> export, with a basic <b>/live</b> , <b>/ready</b> , and <b>/metrics</b> (Prometheus) health server.	Observability stack ready.



1.4	<b>Release MVP</b>	Release v0.1 "AetherClaw Beta." Remove Go/Cobra bloat and deprecated paths.	<b>Success Metrics:</b> <5 MB RSS on LicheeRV-Nano, <500 ms boot time.
-----	--------------------	---	--

## Phase 2: 10x UX & Retention (Weeks 6–8)

**Objective: Introduce high-value, user-facing features to improve activation, habit, and retention loops.**

Step	Action Item	Details & Goal	Source/Target Metric
2.1	<b>Voice-First Integration</b>	Add <b>Voice-first</b> capability: built-in local <b>Whisper-tiny</b> for transcription and <b>piper-rs</b> for Text-to-Speech (TTS).	Voice command support fully functional.
2.2	<b>Skill Marketplace Stub</b>	Set up the infrastructure for the <b>AetherClaw federation + public skill marketplace</b> to share and reuse <b>HEARTBEAT.md</b> snippets.	Git-based marketplace stub available.
2.3	<b>Habit Loop Features</b>	Implement the <b>Daily "Soul Check-in" summary</b> push and the <b>Streaks + memory milestones</b>	Daily summary active. Streak tracking live.



		system. Add the <b>Auto-spawn “Life Coach” sub-agent</b> every Sunday.	
<b>2.4</b>	<b>AetherHub Backend</b>	Build the <b>AetherHub Pro</b> backend to support hosted relays and premium models, integrating a payment processor (Stripe).	Subscription backend ready for testing.

## Phase 3: Monetization & Scale (Weeks 9–12)

**Objective: Launch revenue streams, expand the community, and establish scalability moats.**

Step	Action Item	Details & Goal	Source/Target Metric
<b>3.1</b>	<b>Monetization Launch</b>	Launch the full <b>Marketplace</b> (30% cut on paid skills) and the <b>Freemium/Pro subscription model</b> (\$5/mo for unlimited cloud + priority support).	<b>Success Metric:</b> 30% of users on Pro within 90 days.
<b>3.2</b>	<b>Strategic Hardware</b>	Execute the affiliate marketing and <b>Hardware</b>	Official hardware bundles



	<b>Bundle</b>	<b>bundle</b> with partners like Sipeed and MaixCAM.	launched.
3.3	<b>Enterprise Strategy</b>	Launch the <b>API metering: gateway-as-service</b> for enterprises and the <b>Enterprise self-hosted + relay</b> options.	First enterprise PoC engaged.
3.4	<b>Community Program</b>	Formalize the <b>Community maintainers program</b> with clear revenue-share incentives.	10 community maintainers onboarded.
3.5	<b>Scaling &amp; Moats</b>	Implement global multi-region edge relays (Fly.io/Cloudflare). Enhance the proprietary moat with <b>WASM plugins</b> for sandboxed skills.	Global relay network live. WASM plugin standard finalized.



# Ebook/Whitepaper Version: Public Strategy Overview

## 1. The Problem: Why Edge AI Needs a New Standard

### The Current State

Friction and bloat in existing "Claw-family" forks (PicoClaw/OpenClaw).

- *Content from:* Section 1: Friction & Bloat (Config hell, memory creep, cloud-only, high barrier).

### The AetherClaw Vision

The "AI that disappears into your hardware" narrative.

## 2. The Solution: AetherClaw's Core Principles

### FEW Foundation

Fast, Efficient, Workhorse.

- *Content from:* Section 2: Foundation = FEW (300ms boot, <5MB RSS, observability).



## Local-First, Privacy-First

Hybrid local/cloud routing and file-based personality system.

- *Content from:* Section 3: Performance 10x, Section 5: Proprietary Moat (Hybrid router, file-based personality).

## 3. Value and Adoption Strategy

### Value Proposition

"Second life for decade-old devices," personal always-on agent on \$5–10 hardware.

- *Content from:* Section 1: Value Delivery, Section 4: Desirability.

### Frictionless UX

Overview of TUI wizard, optional web dashboard, and voice-first activation.

- *Content from:* Section 3: UX/UI – Frictionless Activation.

### Retention & Community

AetherClaw network, daily check-ins, streaks, and the skill marketplace.

- *Content from:* Section 3: Retention & Habit Loops.

## 4. Business & Viability

### Business Model

Overview of the **Open-Core Strategy** (MIT core) and **Monetization Levers** (Freemium, Pro subscription, Marketplace cut, Enterprise API).



→ *Content from:* Section 4: Viability, Section 5: Open-Core Strategy.

## Competitive Moat

Highlighting the **Ultra-sandbox**, **WASM skills**, and **Hardware co-design hooks**.

→ *Content from:* Section 5: Proprietary Moat.

## Call to Action:

A summary of the immediate execution command and an invitation to join the community.



# FAQs: AetherClaw Project

## ★ What are the success metrics for the AetherClaw Project's 90-day MVP, as outlined in the roadmap?

The success metrics are:

- 10,000 stars on GitHub in 30 days.
- <5 MB RSS on LicheeRV-Nano hardware.
- 30% of users on the Pro subscription within 90 days.
- 5x lower activation time compared to PicoClaw (measured by init-to-first-tool-call).

## ★ What are the four main phases of the AetherClaw Project's step-by-step execution guide and their respective timeframes?

The four main phases are:

- **Phase 0 – Foundation:** (Week 1–2)
- **Phase 1 – Parity & Performance:** (Weeks 3–5)
- **Phase 2 – 10× UX & Retention:** (Weeks 6–8)
- **Phase 3 – Monetization & Scale:** (Weeks 9–12)

## ★ Which features are included in Phase 2: 10× UX & Retention?

Phase 2 focuses on high-value, user-facing features, including:

- **Voice-First Integration:** Built-in local Whisper-tiny for transcription and Text-to-Speech (TTS) via piper-rs.
- **Skill Marketplace Stub:** Setting up the infrastructure for the public skill marketplace and AetherClaw federation.
- **Habit Loop Features:** Daily "Soul Check-in" summary, streaks and memory milestones, and the auto-spawn "Life Coach" sub-agent every Sunday.
- **AetherHub Backend:** Building the AetherHub Pro subscription backend with Stripe integration.

